

CS771A Mini-Project 2: Continual LwP Model Generation on CIFAR-10 Dataset.

Anirban Sarkar
241110009

Nayan Das
241110045

Shekhar Pratap Singh
241110066

Sushant Gangurde
241110073

Junth Basnet
241110090

Problem Description

Task 1 Description

Task 1 involves training a sequence of models f_1, f_2, \dots, f_{10} using the labeled dataset D_1 and progressively updating them with predicted labels from unlabeled datasets D_2, D_3, \dots, D_{10} . Each model is evaluated on held-out labeled datasets $\hat{D}_1, \hat{D}_2, \dots, \hat{D}_{10}$. The goal is to maintain high accuracy on the current dataset and prevent significant performance degradation on previous datasets.

Task 2 Description

Task 2 involves incrementally training a series of models $f_{11}, f_{12}, \dots, f_{20}$ using datasets $D_{11}, D_{12}, \dots, D_{20}$. Each model is evaluated on its corresponding held-out dataset as well as all previous held-out datasets. The results are presented in a matrix format to compare the performance across different datasets.

Approach Description

Feature Extraction

We employed a feature extraction process using a pre-trained **ConvNeXt Tiny** model. The model, pre-trained on ImageNet, provides robust feature representations that generalize well across similar datasets. The feature extractor uses the penultimate layer of the network to generate meaningful representations of the input images, reducing them to compact and informative vectors.

Steps:

1. Input images from the datasets are resized to 224×224 pixels using OpenCV to match the input size requirement of the pre-trained model.
2. The resized images are passed through the ConvNeXt Tiny model, and outputs from the penultimate layer are extracted.
3. This process converts raw image data into high-dimensional feature vectors, which are stored for subsequent processing.

Raw datasets (D_1 to D_{20}) were initially in PyTorch tensor format. They were converted to NumPy arrays and augmented using the ConvNeXt-based feature extractor. Each dataset is saved as a `.npz` file in augmented dataset folder containing:

- **data**: Feature vectors extracted from the ConvNeXt Tiny model.
- **targets** (if available): Ground truth labels.

LwP Implementation

We implemented the Learning With Prototype (LWP) model, a simple classifier that leverages class prototypes. For each class, the model maintains a prototype vector, which is updated incrementally during training using a weighted moving average. This method combines the existing prototype with the new training sample in proportion to their respective contributions: the previous prototype is weighted by the number of samples already seen for that class, while the new sample is weighted inversely by the total number of samples for the class. This ensures that the prototype evolves as more data is observed, reflecting the overall distribution of the class. By updating prototypes incrementally, the model achieves efficient continual learning without storing past samples. During prediction, it computes distances (or similarities) between the input samples and the prototypes of all classes. The class with the closest prototype is predicted.




Figure 1: t-SNE plot on D_1 dataset

Continual Model Generation

Task 1

Initial Training: A custom LwP (Learning with Prototypes) classifier f_1 is trained on D_1 using the extracted features and their corresponding labels.

Model Update Mechanism: For each unlabeled dataset D_i ($i \geq 2$):

1. Predictions are generated for the dataset using the current model f_{i-1} .
2. Predictions with high confidence (above a threshold, e.g., 0.7) are selected as pseudo-labels.
3. These pseudo-labeled examples are used to update the model, creating f_i .

This ensures that new data enhances the model while minimizing noisy updates from low-confidence predictions. All models f_1 to f_{10} have a consistent architecture and parameter count, ensuring comparability and scalability.

Task 2

Initial Training: Starting with f_{10} as the initial model that was saved as f10_model.pkl file, we proceed with incremental updates to learn models $f_{11}, f_{12}, \dots, f_{20}$. The datasets $D_{11}, D_{12}, \dots, D_{20}$ have slightly different input distributions, requiring an adaptation of the update mechanism from Task 1.

Model Update Mechanism: For each dataset D_i ($i \geq 11$):

1. Predictions are generated for the dataset D_i using the current model f_{i-1} .
2. Predictions with high confidence (above a threshold, e.g., 0.99999) are selected as pseudo-labels.
3. These pseudo-labeled examples are used to update the model, creating f_i , while also considering dataset distribution differences.
 - A domain adaptation technique is applied to address potential shifts in data distributions between D_i and D_{i-1} .

This ensures that new data improves the model while maintaining stability across varying input distributions. Models f_{11} to f_{20} have a consistent architecture and parameter count.

Evaluation Results

Evaluation of Task 1

Evaluation Datasets: Each model f_i is evaluated on its corresponding held-out dataset \hat{D}_i and all preceding held-out datasets $\hat{D}_1, \hat{D}_2, \dots, \hat{D}_{i-1}$.

Accuracy Matrix: A 10×10 accuracy matrix is generated to record the performance of f_i on \hat{D}_j ($1 \leq j \leq i$):

- Rows represent models f_1, f_2, \dots, f_{10} .
- Columns represent datasets $\hat{D}_1, \hat{D}_2, \dots, \hat{D}_{10}$.

	D1	D2	D3	D4	D5	D6	D7	D8	D9	D10
f1	0.8732									
f2	0.8648	0.8792								
f3	0.8636	0.878	0.8704							
f4	0.8604	0.8772	0.8684	0.872						
f5	0.8596	0.876	0.8688	0.8728	0.8712					
f6	0.8580	0.8764	0.8684	0.8728	0.87	0.8624				
f7	0.8564	0.8756	0.8688	0.872	0.8716	0.8628	0.86			
f8	0.8572	0.8748	0.8684	0.8692	0.8704	0.8616	0.8584	0.8692		
f9	0.8568	0.8744	0.8684	0.8704	0.8712	0.8608	0.8584	0.8684	0.8692	
f10	0.8560	0.8752	0.8688	0.8688	0.8708	0.8596	0.8584	0.8684	0.8704	0.8764

Figure 2: Evaluation Matrix of Task 1

The approach ensures minimal performance degradation on earlier datasets as the models are sequentially updated.

Evaluation of Task 2

Accuracy Matrix: A 10×20 accuracy matrix is generated to record the performance of models $f_{11}, f_{12}, \dots, f_{20}$ on held-out datasets $\hat{D}_1, \hat{D}_2, \dots, \hat{D}_{20}$:

- Rows represent models $f_{11}, f_{12}, \dots, f_{20}$.
- Columns represent held-out datasets $\hat{D}_1, \hat{D}_2, \dots, \hat{D}_{20}$.

This matrix highlights the trade-off between adapting to new datasets and retaining performance on earlier datasets.

	D1	D2	D3	D4	D5	D6	D7	D8	D9	D10	D11	D12	D13	D14	D15	D16	D17	D18	D19	D20
f11	0.8556	0.8744	0.8672	0.8680	0.8692	0.8580	0.8568	0.8676	0.8684	0.8740	0.7332									
f12	0.8528	0.8744	0.8668	0.8672	0.8684	0.8560	0.8580	0.8660	0.8672	0.8728	0.7332	0.512								
f13	0.8528	0.8732	0.8660	0.8660	0.8684	0.8552	0.8576	0.8664	0.8660	0.8724	0.7312	0.5112	0.7784							
f14	0.8512	0.8728	0.8656	0.8648	0.8688	0.8560	0.8560	0.8656	0.8656	0.8720	0.7300	0.5104	0.7788	0.8232						
f15	0.8516	0.8720	0.8660	0.8656	0.8668	0.8552	0.8540	0.8652	0.8636	0.8704	0.7312	0.5116	0.7772	0.8228	0.8504					
f16	0.8492	0.8732	0.8664	0.8644	0.8664	0.8552	0.8532	0.8640	0.8636	0.8696	0.7344	0.5128	0.7756	0.8216	0.8508	0.7236				
f17	0.8500	0.8728	0.8652	0.8640	0.8668	0.8552	0.8532	0.8640	0.8632	0.8688	0.7336	0.5112	0.7756	0.8216	0.8508	0.7248	0.7736			
f18	0.8492	0.8720	0.8636	0.8628	0.8660	0.8552	0.8528	0.8632	0.8628	0.8688	0.7320	0.51	0.774	0.8224	0.8492	0.7232	0.7724	0.7416		
f19	0.8484	0.8724	0.8648	0.8620	0.8664	0.8536	0.8536	0.8620	0.8612	0.8680	0.7308	0.5104	0.774	0.8212	0.8484	0.7224	0.7724	0.7388	0.6496	
f20	0.8492	0.8708	0.8640	0.8628	0.8648	0.8524	0.8540	0.8608	0.8624	0.8672	0.7320	0.5116	0.7736	0.8216	0.8464	0.7224	0.7716	0.7384	0.6468	0.8284

Figure 3: Evaluation Matrix for Task 2

Performance Analysis: The accuracy matrix is analyzed to identify trends and patterns in model performance.

1. After training f_i on D_i , evaluate f_i on the corresponding held-out dataset \hat{D}_i .
2. Additionally, evaluate f_i on all preceding held-out datasets $\hat{D}_1, \hat{D}_2, \dots, \hat{D}_{i-1}$.
3. Compare the performance of f_i on earlier held-out datasets against f_{i-1} to ensure minimal degradation.

Conclusion

Based on the accuracy matrices from both Task 1 and Task 2, we can draw several conclusions about the performance and progression of the models:

Task 1 Summary

- **Model Progression:** The models f_1 through f_{10} were trained sequentially on datasets D_1 to D_{10} . Each model was evaluated on its corresponding held-out dataset as well as all previous datasets.
- **Accuracy Trends:** The accuracy generally decreases as the model is applied to datasets it was not directly trained on, which is expected due to the potential differences in data distribution.
- **Final Model Performance:** The final model, f_{10} , achieved a reasonable accuracy across all datasets, indicating that the model was able to generalize well from the sequential training process.

Task 2 Summary

- **Model Progression:** Starting with f_{10} from Task 1, models f_{11} through f_{20} were trained on datasets D_{11} to D_{20} . Each model was evaluated on its corresponding held-out dataset as well as all previous datasets.
- **Accuracy Trends:** Similar to Task 1, the accuracy tends to decrease when models are evaluated on datasets they were not directly trained on. However, the models show a consistent ability to maintain reasonable accuracy across multiple datasets.
- **Final Model Performance:** The final model, f_{20} , shows a good level of accuracy across the board, suggesting that the sequential training approach continues to be effective in Task 2.

Overall Conclusion

- **Sequential Training Effectiveness:** The sequential training approach used in both tasks appears to be effective in building models that can generalize across multiple datasets. This is evidenced by the relatively high accuracy scores across both tasks.
- **Model Generalization:** The models demonstrate a strong ability to generalize, as seen by their performance on datasets they were not directly trained on.
- **Potential for Improvement:** While the models perform well, there is always room for improvement, particularly in handling datasets with significantly different distributions or characteristics.

These results suggest that the methodology used is robust and can be applied to similar tasks where sequential learning and evaluation are required.

Problem 2

Liu, Chenxi, et al. "DEJA VU: Continual Model Generalization For Unseen Domains." [Online]. Available: <https://arxiv.org/abs/2301.10418>, 2023.

Youtube link: <https://www.youtube.com/watch?v=-QMVvgShvB0>

Acknowledgments

We would like to express our deepest gratitude to Professor Piyush Rai for designing this insightful CS771A Machine Learning course and for providing clear guidance and thought-provoking challenges throughout this mini-project. His expertise in the field of machine learning, coupled with his comprehensive approach to teaching, has been instrumental in our learning journey. We also thank the teaching assistants for their support in answering questions and clarifying doubts during the course and the project. We acknowledge the importance of the tools and frameworks that were integral to our project, including Scikit-learn, Keras, and TensorFlow, for their role in helping us implement machine learning models effectively.

References

- Liu, Zhuang, et al. "A ConvNet for the 2020s." *CoRR*, Volume abs/2201.03545, 2022. [Online]. Available: <https://arxiv.org/abs/2201.03545>
- Liu, Chenxi, et al. "DEJA VU: Continual Model Generalization For Unseen Domains." [Online]. Available: <https://arxiv.org/abs/2301.10418>, 2023.
- He, Kaiming, et al. "Deep Residual Learning for Image Recognition." [Online]. Available: <https://arxiv.org/abs/1512.03385>, 2015.
- Chollet, François, et al. "Keras: Deep Learning library for Theano and TensorFlow." 2015. [Online]. Available: <https://keras.io>
- Chollet, François. *Deep Learning with Python.* 2nd ed., Manning Publications, 2021.
- Géron, Aurélien. *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems.* 3rd ed., O'Reilly Media, 2022.
- Stanford CS229: Machine Learning Course — Summer 2019 (Anand Avati). Stanford Online, [YouTube Playlist]. Available: https://www.youtube.com/watch?v=KzH1ovd40ts&list=PLoROMvovd4rNH7qL6-efu_q2_bPuy0adh
- MIT 6.S191: Convolutional Neural Networks. MIT OpenCourseWare, [YouTube Video]. Available: <https://www.youtube.com/watch?v=2xqkSUhmmXU>
- Deep Learning With Tensorflow 2.0, Keras. [YouTube Playlist]. Available: https://www.youtube.com/playlist?list=PLeo1K3hjS3uu7CxAacxVndI4bE_o3BDt0